

# AI 0→1 Learning Path

Purpose: a practical, timeline-free path. For each step,

**Learn the items**

**Tick the checklist**

---

## Step 1 — Python Core

**Learn:** syntax, control flow, functions & modules, error handling; data structures (list/tuple/set/dict), list/dict comprehensions; basic CLI scripting.

### Checklist

- ☐ Write 10 small exercises covering loops, functions, slicing, and exceptions.
- ☐ Implement a CSV→JSON CLI with argparse and basic logging.
- ☐ Explain mutability, shallow vs deep copy, and when to use each container.
- ☐ Set up a virtual environment and install a third-party package.

---

## Step 2 — Intermediate Python

**Learn:** file I/O for CSV/JSON/XML; OOP (classes, inheritance, dataclasses), typing; iterators/generators/decorators; context managers; packaging and environment management (uv/poetry or venv+pip).

### Checklist

- ☐ Build a small I/O utility library (read/write CSV, JSON, XML) with docstrings.
- ☐ Add unit tests with pytest (cover ≥80%).
- ☐ Add pre-commit with ruff/black/isort and ensure it runs locally.
- ☐ Publish the library in a private GitHub repo with a clear README.

---

## Step 3 — Scientific Computing & Data Stack

**Learn:** NumPy (arrays, broadcasting, vectorization), Pandas (indexing, groupby, merge, missing values), Matplotlib (core plots; Seaborn optional), Data Quality checks.

### Checklist

- ☐ Produce an EDA notebook on a public dataset with at least 3 plots and key findings.
  - ☐ Deliver a Data Quality Report (source, schema, missingness, outliers, assumptions).
  - ☐ Demonstrate vectorized NumPy operations vs Python loops with timing.
- 

## Step 4 — Math Foundations (with Python)

**Learn:** Linear Algebra (norms, eigen/singular values intuition), Calculus/Optimization (derivatives, gradients, LR schedules), Probability & Statistics (distributions, sampling, estimation, confidence intervals, bias–variance trade-off).

### Checklist

- ☐ Implement gradient descent from scratch for linear regression; compare to sklearn.
  - ☐ Plot learning curves and diagnose under/overfitting.
  - ☐ Explain in your own words: variance vs bias, regularization, and early stopping.
- 

## Step 5 — Classical Machine Learning (scikit-learn)

**Learn:** Supervised (linear/logistic, trees/forests, gradient boosting), unsupervised (k-means, PCA); feature engineering (scaling/encoding/selection); model evaluation (train/val/test, CV, metrics, leakage); Pipeline & ColumnTransformer.

### Checklist

- ☐ Build an end-to-end sklearn Pipeline (preprocess→model→evaluate) with CV.

- ☐ Compare at least 3 models and report metrics (ROC/PR/F1) and feature importance.
  - ☐ Provide a model card and data card; ensure a fresh clone can reproduce your score.
- 

## Step 6 — Deep Learning Basics

**Learn:** choose PyTorch or TensorFlow; tensors & autograd; optimizers (SGD/Adam); regularization (dropout/weight decay); scheduling; early stopping; CNNs for images; saving/loading checkpoints.

### Checklist

- ☐ Train a small CNN on CIFAR-10 (or similar) and achieve a baseline accuracy.
  - ☐ Save best checkpoints; plot loss/accuracy curves; create a confusion matrix.
  - ☐ Explain vanishing/exploding gradients and remedies (init, normalization, skip connections).
- 

## Step 7 — Advanced Topics (modular, pick and combine)

### 7A. LLM & RAG (recommended)

**Learn:** chunking strategies, embeddings, vector DBs (FAISS/pgvector), retrieval→reranking→answering pipelines, evaluation (exact match, semantic similarity, citation hit rate, hallucination rate), lightweight fine-tuning (LoRA/QLoRA), inference optimization (quantization, batching, caching).

### Checklist

- ☐ Build a PDF Q&A assistant (docs→chunks→embeddings→retrieval→answer) with an API.
- ☐ Create an evaluation set; report precision@k, citation hit rate, hallucination rate.
- ☐ Run an A/B of prompt templates or rerankers; document latency/cost trade-offs.

## 7B. NLP / CV / RL (choose one to deepen)

### Learn:

- NLP: tokenization, classic features vs transformers, sequence classification.
- CV: OpenCV basics, modern CNN/ViT intuition, augmentations.
- RL: agents, environments, policies, reward design.

### Checklist

- ☐ Deliver one small demo (e.g., text classifier, image classifier, or basic RL agent).
- ☐ Include error analysis and at least one robustness test.

## 7C. MLOps (strongly encouraged)

**Learn:** Docker, FastAPI inference service, experiment tracking (MLflow or W&B), data/model versioning, basic monitoring (latency, throughput, failure rate, drift alerts), CI with GitHub Actions.

### Checklist

- ☐ Containerize training and inference; expose a /predict endpoint.
  - ☐ Track experiments (params, metrics, artifacts) and keep runs reproducible.
  - ☐ Add CI to lint/test; provide a Makefile for make setup/train/eval/serve.
- 

## Step 8 — Engineering Best Practices

**Learn:** Git branching; code style (PEP8 with ruff/black/isort), typing (mypy/pyright), testing pyramid (unit→integration), documentation (docstrings/README), observability (structured logs/metrics), basic performance testing.

### Checklist

- ☐ New teammate can clone and run make setup && make train && make serve without edits.
- ☐ Linting and tests pass locally and in CI; type checks clean.
- ☐ Architecture and data-flow diagram included in README.

---

## Self-Check Rubrics (tick when you can convincingly demo)

- Reproducibility: fresh environment can re-train, re-evaluate, and start the API.
  - Metrics literacy: can choose metrics, explain AUC vs PR, avoid leakage.
  - Training diagnostics: can address overfitting with concrete interventions.
  - RAG quality: can measure and reduce hallucinations; provide failure cases.
  - Ops readiness: images build, P95 latency within target, basic alerts wired.
- 

## Project Skeleton

```
project-name/  
├── data/ (raw, interim, processed)  
├── notebooks/  
├── src/ (features, models, training, inference)  
├── tests/  
├── configs/ (yaml)  
├── deployment/ (Dockerfile, compose.yaml)  
├── Makefile  
├── pyproject.toml (or uv/poetry)  
└── README.md (model/data cards)
```

---

## Resources

### Official Documentation (primary)

- scikit-learn User Guide — [https://scikit-learn.org/stable/user\\_guide.html](https://scikit-learn.org/stable/user_guide.html)
- PyTorch Tutorials — <https://pytorch.org/tutorials/>
- TensorFlow Learn/Guide — <https://www.tensorflow.org/learn>
- Hugging Face LLM Course — <https://huggingface.co/learn/llm-course>

- MLflow Documentation — <https://mlflow.org/docs/>
- Weights & Biases Docs — <https://docs.wandb.ai/>
- FAISS Docs — <https://faiss.ai/>
- Kaggle Learn — <https://www.kaggle.com/learn>

## **Blogs (read as needed)**

- Machine Learning Mastery — <https://machinelearningmastery.com/>
- fast.ai Blog — <https://www.fast.ai/posts/>
- Google AI Blog — <https://ai.googleblog.com/>
- OpenAI Blog — <https://openai.com/research>
- Towards Data Science — <https://towardsdatascience.com/>
- KDnuggets — <https://www.kdnuggets.com/>
- Analytics Vidhya — <https://www.analyticsvidhya.com/blog/>
- Real Python — <https://realpython.com/>
- DataCamp Tutorials — <https://www.datacamp.com/tutorial>

## **YouTube Channels (pick a few, not all)**

- freeCodeCamp.org — <https://www.youtube.com/@freecodecamp>
- Corey Schafer — <https://www.youtube.com/@coreyms>
- Tech With Tim — <https://www.youtube.com/@TechWithTim>
- StatQuest with Josh Starmer — <https://www.youtube.com/@statquest>
- Sentdex — <https://www.youtube.com/@sentdex>
- Data School — <https://www.youtube.com/@dataschool>
- DeepLearningAI (Andrew Ng) — <https://www.youtube.com/@Deeplearningai>
- Two Minute Papers — <https://www.youtube.com/@TwoMinutePapers>
- Krish Naik — <https://www.youtube.com/@krishnaik06>
- Codebasics — <https://www.youtube.com/@codebasics>

- CS50 AI (Harvard) — <https://www.youtube.com/playlist?list=PLhQjrBD2T382Nz7z1AEXmioc27axa19Kv>
-